
The Elicita Handbook

Elicita helps you ask the question you would otherwise have missed — while the client is still in the room — and then writes up what was agreed, quoting the words that were actually said. This is a guide to using it. It goes in the order you will meet it: setting it up, preparing for an engagement, the meeting itself, and what comes out afterwards. Every screen is shown. It assumes nothing about how the software is built. Where something is not finished, it says so plainly rather than describing an intention in the same voice as a fact.

Contents

GETTING STARTED

What Elicita Does	3
Setting It Up	5
Getting It Onto People's Machines	10

BEFORE A MEETING

Preparing for an Engagement	12
Starting a Meeting	14

DURING THE MEETING

The Panel and the Suggestion	17
Controlling the Recording	20
Running a Meeting in Two Languages	23
When the Connection Drops	25

AFTER THE MEETING

Checking the Recording	27
Getting the Write-Up	29
Carrying It Into the Next Meeting	33

RUNNING IT WELL

Judging Whether the Suggestions Are Any Good	35
What It Will Not Do	38
What Works Today	39

GETTING STARTED


What Elicita Does

Imagine a meeting where a client is describing what they want built. Somebody says *"we need the reports to be fast."*

That sounds like a requirement. It is not one. Fast for whom? Compared to what they have now? A second, a minute, overnight? You almost certainly know to ask — and quite often do not, because at that moment you are also listening, taking notes, watching the clock, keeping the room comfortable, and trying to remember which topics you have not covered yet.

So the moment passes. "Fast" goes into the notes as though it meant something. Months later somebody builds a report that takes four seconds, the client says that is not what they asked for, and it gets built again.

Elicita watches for that moment and puts the question in front of you while there is still time to ask it.



A suggestion, with the short version you can read at a glance and the reason it fired underneath

The shape of it

What you do, and what Elicita does

1

Days before

You tell it about the client and give it what you already have. It drafts the questions worth asking. You read them, reorder them, and delete the ones that miss.

2

In the room

It listens. When something is left vague or a topic gets skipped, one suggested question appears. You decide whether to ask it, with a single tap.

3

Afterwards

It writes up what was agreed — the open questions, the decisions, a project brief, a follow-up email — with the words that were actually said attached to every claim.

The thinking is done in the first stage. The meeting only picks from work already finished.

What it is like to use

It is quiet. A good meeting might produce three suggestions, not twelve. When it is unsure, it says nothing, because one suggestion that embarrasses you in front of a client costs far more than a good one earns.

Every response is a single tap. There is nothing to type, nothing to read through, and no conversation to have with it. It offers; you accept or you do not.

And it never speaks to the client. It has no voice in the meeting, sends nothing on your behalf, and joins nothing as a participant that talks. If you ignore it for ninety minutes, nothing happens — that is a supported way to use it.

Who this guide is for

The person running the meeting, and whoever sets Elicita up for them. It covers what each screen does and how to use it, in order: setting it up, preparing for an engagement, the meeting itself, and what comes out afterwards.


It assumes nothing about how the software is built, and nothing about the company being interviewed. Where something is not finished yet, this guide says so rather than describing what is intended as though it already worked.

Setting It Up

Elicita uses outside services for two things: turning speech into text, and doing the writing-up afterwards. Somebody has to tell it which services to use and give it permission to use them. It is done in the app, once, by whoever administers Elicita for the team — not by an engineer editing a server somewhere, because the person who needs a key changed is rarely the person who could deploy one.

Before anything is set up

On first run the screen says plainly that nothing is configured, rather than looking ready and failing later.




The screenshot shows the 'Settings' screen of the Elicita application. At the top, there's a 'SERVICE' header with a circular icon. Below it, the title 'Settings' is displayed. A warning box states: 'These settings are held in memory and will be lost when the service restarts.' The 'AI PROVIDER' section is titled 'Route Claude calls through' and shows 'Anthropic' selected in a dropdown menu. The 'CREDENTIALS' section has a sub-header 'Authenticate with' and two buttons: 'API key' and 'OAuth token'. Below this, there are two configuration blocks. The first is for 'Anthropic API key', which shows 'Not configured' and a 'Test' button. The second is for 'Anthropic OAuth token', also showing 'Not configured' and a 'Test' button. Both blocks include a note: 'No key is stored yet.'

Before anything is configured

Choosing where the writing-up happens

You can point Elicita at the company that makes the underlying service directly, at Amazon's, Google's or Microsoft's hosted versions of it, or at your organisation's own internal gateway. Each choice asks only for what it actually needs, and refuses to save if something is missing — a message here is better than a failure halfway through writing up a meeting.

Two of those choices ask for no key at all. If you already run on Amazon or Google, Elicita uses the permissions your cloud account grants it. Asking you to paste a key there would mean creating a long-lived credential in a place where your own platform has a better mechanism.



The screenshot shows the 'Settings' page for a 'SERVICE'. Under the 'AI PROVIDER' section, the 'Route Claude calls through' dropdown is set to 'Other Anthropic-compatible endpoint'. The 'Endpoint' field contains 'https://llm.internal/v1'. Under the 'CREDENTIALS' section, the 'Authenticate with' radio buttons have 'API key' selected. The 'Anthropic API key' field shows 'Configured — ends x7q2' with a green 'in use' badge, and 'Test' and 'Clear' buttons. The 'Anthropic OAuth token' field is 'Not configured' with a 'Test' button.

Pointing Elicita at an internal gateway

Keys and tokens

Both are supported, because organisations issue different things. You choose which is in use and Elicita labels it, so there is never any doubt about which credential a failure came from.

The screenshot shows the 'Settings' page for a 'SERVICE'. It is divided into two main sections: 'AI PROVIDER' and 'CREDENTIALS'.

AI PROVIDER

Route Claude calls through
Calls api.anthropic.com with the credential below.
A dropdown menu shows 'Anthropic' selected.

CREDENTIALS

Authenticate with
Bring your own credential. Use whichever your organisation issues — a key or a token, not both.
Two buttons are present: 'API key' (selected) and 'OAuth token'.

Anthropic API key in use
A key issued from the Anthropic console (starts sk-ant-).
A text field shows 'Configured — ends x7q2'. To its right are 'Test' and 'Clear' buttons.
Below the field, it says: 'A key is stored. Leave this blank to keep it.'

Anthropic OAuth token
A token from 'claude setup-token', if your organisation issues those instead of keys.
A text field shows 'Not configured'. To its right is a 'Test' button.
Below the field, it says: 'No key is stored yet.'

Speech-to-text vendor key
Used to transcribe the meeting on both the live and record paths.
A text field shows 'Configured — ends k3m8'. To its right are 'Test' and 'Clear' buttons.

Credentials configured, with the one in use marked

Four details on this screen exist to stop a credential leaking, and they are worth knowing about because they change how the screen behaves:

Why the credential fields behave oddly

The field is always empty

Elicita never sends a stored key back to the screen, so there is nothing to show. It displays the last four characters instead, which is enough to tell one key from another.

Blank means leave it alone

Saving a different setting cannot wipe your key by accident.

Clear is separate and deliberate

Removing a credential is its own action, because it cannot be undone.

Test really tests

It checks the credential against the service rather than reporting "configured" as though that meant "working". If it fails you see the service's own explanation, never your key.

Choosing transcription services

The live transcription and the after-the-meeting transcription are set up separately, because they are bought on different qualities. The live one is chosen for how quickly it can tell that a sentence has finished. The after-the-meeting ones are chosen for being accurate — and for disagreeing usefully with each other.

Managed capture vendor key
 Used to join the meeting and capture per-participant audio.

Not configured
Test

 No key is stored yet.

INFERENCE

Model
 Used for the context compiler and the debrief pipeline.
 claude-opus-5

SPEECH VENDORS

Live transcription
 Drives the in-meeting nudges. Ends a turn when the sentence sounds finished, which cuts the wait before a nudge.
 AssemblyAI

Recording transcription
 Two engines transcribe the recording separately after the meeting; where they disagree is flagged for you to check. They must be different vendors — the same engine twice would always agree with itself.
 Deepgram + AssemblyAI

☒ **Send engagement vocabulary to the transcriber**
 Client and product names are transcribed far more accurately when the engine is told about them in advance.

☒ **Tell vendors not to retain client audio**
 Sets the opt-out on every request. Your contract may already say this; sending it per request is what an audit can verify.

Transcription services

The two used for the recording must be different companies. The same service twice would agree with itself and flag nothing, so the form will not accept it. You can also point Elicita at a service it does not ship support for, though it will tell you that some behaviour then needs checking against that provider yourself.

For a shared installation


Where several machines share a deployment, the key can come from your own secret store rather than being typed in — anything with a command line will do. If the store cannot be reached, Elicta refuses to start and says why, rather than quietly generating a new key that would leave the other machines unable to read anything.

Getting It Onto People's Machines

This is invisible when it goes well and very visible when it does not. A meeting tool that asks every person for permission to record their screen generates the same support ticket over and over.


The preferred route

Push Elicita out through your device management system. That is the recommended route rather than the fallback, because it lets you grant microphone and screen-recording access centrally. Nobody sees a permission prompt, and nobody has to be talked through one during a client meeting.



A managed installation, showing what was granted centrally

Keep the plain installers as a backup for machines outside management.



The screenshot displays the Elicta application interface. At the top left, there is a circular logo and the word "ABOUT". Below this, the Elicta logo is shown next to the text "Elicta" and "0.1.0 (build 274) · Windows 11 x64".

The "THIS BUILD" section contains three items:

- Signed and notarised**: Developer ID — Northwind Consulting Pty Ltd — the OS verified this before it ran. A green "Verified" badge is present.
- Installed by Direct download**: Direct download. Permissions were granted by you at first launch.
- Updates**: Checks weekly.

The "PERMISSIONS" section contains two items:

- Microphone**: Captures the meeting audio. A green "Granted" badge is present.
- Screen Recording**: Required for loopback capture of a silent join. A red "Needed" badge is present.

At the bottom of the permissions section, a note states: "Capture will not start without Screen Recording. On a managed machine your IT team can grant this by profile rather than each person answering a prompt."

An unmanaged installation

Agree the security exception first

An unfamiliar application opening an audio stream is exactly what endpoint security software is built to flag. It will flag it — and if this is left until the pilot, it will flag it in the middle of somebody's client meeting.

Get the exception agreed before the pilot rather than during it. This is the single most common cause of a bad first impression.

What each person does at first launch

Very little, if the push worked. Otherwise they will be asked once for microphone access and, on some systems, for screen recording. Both are required for Elicta to hear the meeting at all.

BEFORE A MEETING

Preparing for an Engagement

Almost everything that makes Elicita useful happens before anyone sits down. That is deliberate: the thinking is done in advance so that during the meeting it only has to pick the right question from work already finished. It is the only way a suggestion can arrive while it is still worth having.

Do this days before the first meeting, not the morning of it.

Tell it about the client

Who they are, the sector, and the commercial shape of the engagement.

That is not filing. It is what lets Elicita work out which languages to expect in the room, so that nobody has to be asked to choose one before a meeting starts.

Add what you already have

Scoping decks, throughput studies, the last proposal. Each document is tagged as one of three things, and the tag changes how Elicita treats it:

The three document tags

Ground truth	Established fact. If a client contradicts it, that is worth interrupting for.
Hypothesis	Our assumption, not theirs. It produces questions to verify rather than contradictions to flag.
Superseded	Kept for background, never used to challenge anyone.

The tag decides whether a contradiction is worth interrupting a meeting for.

Add the client's own words

Product names, internal systems, their own jargon — the words a transcription service would have no way of knowing.

This is the single most effective thing you can do for accuracy, and it is worth more than any other preparation on this page. A misheard product name does not just look sloppy. It reads to Elicita as something brand new, and it will interrupt you to ask about a system that does not exist.

Review the questions before anybody asks them

Elicita reads everything and drafts the questions worth asking, grouped by the section of your requirements template that each one serves. You read them, reorder them, and delete the ones that miss.

ENGAGEMENT

Northwind Logistics

REFERENCE DOCUMENTS

Current-state scoping deck Indexed for retrieval	ground truth
Warehouse throughput study 2025 Indexed for retrieval	ground truth
Proposed integration approach Indexed for retrieval	hypothesis
Original RFP response Indexed for retrieval	superseded

A document tagged **ground truth** is the only kind a contradiction can fire against.

ENGAGEMENT VOCABULARY

Northwind Zephyr WMS Consignment Cross-dock SLA-4

Sent to the transcriber as keyterms. Client and product names are the words most often misheard, and a misheard product name fires a trigger about nothing.

QUESTION BANK

Performance 2

When you say the dashboard has to be fast, what does that mean in seconds? Priority 1	Prune
Is that response time at median load, or at your Monday morning peak? Priority 2	Prune

Integrations 2

Which systems does Zephyr WMS have to talk to on day one? Priority 1	Prune
Are those integrations real-time, or is a nightly batch acceptable? Priority 2	Prune

The preparation screen: documents, the client's vocabulary, and the question bank grouped by section

This review is worth doing even if you never open Elicita during the meeting. A good question list makes you better prepared on its own — which is why it comes first, and why you are meant to edit it rather than trust it.

One thing that needs setting up first

Drafting the questions needs one of the outside services configured. Until that is done the screen works and the question bank stays empty.

Starting a Meeting

Recording a client without a clear record of their consent is the kind of mistake that ends an engagement. So in Elicita it is a gate you pass, not a warning you read past.

Consent first

The screen shows how consent works for this client — agreed once for the whole engagement, or needed for every meeting — and whether it has actually been given.

Until somebody confirms it on the record, recording cannot start. The button is simply unavailable.

 **BEFORE RECORDING**

Discovery 3 — integrations and volumes

CONSENT

Consent model
per meeting

Not confirmed
Capture cannot start until someone confirms this on the record.

Required

WHAT HAPPENS TO THE AUDIO

Never written to disk
Audio stays in memory for the duration of transcription.

Destroyed after transcription
The moment the recording has been transcribed and diarized, the audio is deleted and the deletion is recorded.

Vendors told not to retain
Sent as a parameter on every request, not only in the contract.

On

CAPTURE

Line-in from the meeting machine
Platform voice processing is off — it is tuned for a listener, and removes detail the transcriber uses.

Start

Consent not yet confirmed, and recording cannot start

Once confirmed, the screen keeps who confirmed it and when. The answer to "did we have permission for this?" is never a memory test.


BEFORE RECORDING

Discovery 3 — integrations and volumes

CONSENT

Consent model
standing for the engagement

Confirmed
Priya Raman (delivery lead) · 18 Aug 2026, 09:58

On record

WHAT HAPPENS TO THE AUDIO

Never written to disk
Audio stays in memory for the duration of transcription.

Destroyed after transcription
The moment the recording has been transcribed and diarized, the audio is deleted and the deletion is recorded.

Vendors told not to retain
Sent as a parameter on every request, not only in the contract.

On

CAPTURE

Line-in from the meeting machine
Platform voice processing is off — it is tuned for a listener, and removes detail the transcriber uses.

Start

Consent confirmed, and recording can begin

What happens to the audio, said before it happens

The same screen tells you what Elicita will do with the recording, before a second of it exists. You can read it aloud to a client who asks.

What happens to the recording

- 1 It is never written to disk**
The audio lives in memory only, and only while it is being turned into text.
- 2 It is destroyed as soon as it is used**
The moment transcription and speaker identification finish, the recording is deleted — and the deletion is itself recorded, so the destruction can be shown rather than asserted.
- 3 The transcription service is told not to keep it**
Every request carries that instruction. Your contract may already say so; sending it with each request is what makes it something an audit can check.

Each of these is something Elicita does, not something it promises.

Joining the meeting

Elicita can take audio from a wired input on the machine, or join the call as a silent participant that never speaks. The silent join is markedly more accurate, for a reason covered in the chapter on controlling the recording: people talking over each other causes more transcription errors than the choice of service does.


DURING THE MEETING

The Panel and the Suggestion

This is what Elicita is for. A client says something that sounds like an answer but is not one, and five seconds later you ask the question that pins it down — instead of thinking of it in the car afterwards.

Before anyone speaks

The panel shows how much of your template is covered and how long is left, and nothing else. There is no suggestion because nobody has said anything yet.




The panel before the meeting starts

An empty panel is the correct resting state. Anything shown here would pull your eyes to a screen during the part of the meeting where you are establishing rapport.

Something vague lands

The client says *"the dashboard has to be fast."* That is not a requirement, it is a feeling. Elicita recognises it and puts one question in front of you.



A suggested question, and why it fired

Three things, in the order you need them:

What is on a suggestion, and why

The short version	Readable at a glance, without breaking eye contact.
The full question	The exact wording, if you want to ask it word for word.
Why it fired	Quietly underneath. This matters more than it looks — it is how you tell in half a second whether Elicita understood the room or misheard it, and therefore whether to trust the next one.

Only one suggestion is shown at a time. Earlier ones move below a line and fade back — still readable if you want them, never competing with the current one.


You answer in one tap

You have one hand and no attention, so every response is a single tap.

The four responses

	What you tap	What it means
You covered that topic	Asked it	The section fills in immediately
Fair question, wrong moment	Park it	It is kept for the write-up
That answer opened something	Go deeper	It follows the thread
Show me the biggest hole	What am I missing?	It names the most urgent uncovered topic

Tapping *Asked it* fills the section in front of you straight away, rather than waiting for anything to confirm. Your confirmation is what you can see.



Coverage moves the moment you tap

The way out

There is also a text box at the bottom for typing your own question. It is deliberately the quietest thing on the screen — a way out, not the way to work.

It exists because a tool that only accepts the four answers it expected will eventually be wrong in a way nobody can tell it about.


Controlling the Recording

Two of the controls on this screen are safety features rather than conveniences, and that shapes everything about how it looks.

Pause is one tap, and always in the same place

You reach for pause when a client says *"can we take this bit off the record."* A control that needs hunting for — or that asks whether you are sure — fails at exactly that moment.

So it runs the full width of the screen, never moves, and takes effect immediately. There is no buffered audio still on its way.



Recording, with pause one tap away

You can tell the state without reading carefully

Believing you are paused while still recording is the worst thing this product could do to you.

So the state is said three ways at once — the word, the sentence under it, and the colour — and never by

colour alone, which also means it works for anyone who cannot distinguish those colours.

CAPTURE

Paused

STATE

Paused 23:41
No audio is reaching the transcriber. Nothing said now is captured.

Resume recording

Pausing takes effect immediately — there is no buffered audio to flush.

INPUT

USB interface — line-in from the meeting machine In use
Wired — the cleanest signal, and the recommended source.

Loopback from a silent join
Loopback from a silent join — clean, and per-participant where the platform allows.

Built-in microphone
Room microphone — picks up echo and cross-talk, which costs accuracy.

YOUR VOICE

Enrolled Re-record
48s sample. Your speech is tagged as yours, so a question you ask is not mistaken for a client requirement.

Paused, and saying so unmistakably

Audio captured while paused is discarded rather than held back, and resuming is instant because the microphone is never actually closed.

It will argue for a better microphone

Elicita does not simply accept whatever audio it is given. A wired input, or joining the call as a silent participant, is markedly more accurate than a microphone in the middle of a table. People talking over each other causes more transcription errors than the choice of transcription service does.

So a room microphone earns a warning that stays on screen, rather than a dialog at the start that everyone dismisses without reading.

 **CAPTURE**

Recording

STATE

Recording 23:41
Audio is being captured and transcribed.

Pause recording

Pausing takes effect immediately — there is no buffered audio to flush.

INPUT

USB interface — line-in from the meeting machine
Wired — the cleanest signal, and the recommended source.

Loopback from a silent join
Loopback from a silent join — clean, and per-participant where the platform allows.

Built-in microphone **In use**
Room microphone — picks up echo and cross-talk, which costs accuracy.

You are on a room microphone. Wired or loopback capture is noticeably more accurate — cross-talk is the single largest source of transcription error.

YOUR VOICE

Not enrolled **Enrol**
Record up to 60 seconds so your own speech can be told apart from the client's.

On a room microphone, with the warning that earns its place

Record your own voice once

You can record up to a minute of yourself so that Elicita can tell your speech from the client's.

It is worth the minute. Without it, a question *you* asked can be recorded as something the client wanted — and that error carries all the way into the write-up, where nobody would know to doubt it.

Running a Meeting in Two Languages

A bilingual meeting is not an English meeting with a few foreign words in it. Requirements arrive in both languages, and the vague answer you need to chase is as likely to be in one as the other.

Both languages, visibly

The panel shows which languages it is hearing, and how well it supports each one.

— — — — 1 / 4 22:00 EN ZH TIER 1

How many, exactly?

You mentioned 三百五十万 orders — is that 3,500,000 a year or a month?

unquantified quantity — numeral grouping

How fast is fast?

unquantified adjective — "fast"

Asked it Park it What am I missing?

Go deeper

Type a question instead...

A suggestion in a meeting running in two languages

That is not decoration. Support varies by language, and you should be able to see at a glance whether a quiet panel means "nothing worth asking" or "I am on unfamiliar ground here."

The suggestion is in the room's language; everything else is in yours

The question you might read aloud appears in the language the meeting is being held in. The short version, the reason it fired, and every label around it stay in yours — and you set that separately.

The person reading the panel and the person being asked the question are not the same person, so they do not get the same language.

Numbers

A spoken figure has to reach you as a figure you can act on: three hundred and fifty myriad has to arrive as **3,500,000**, not as words.

No transcription service handles that reliably, so Elicita converts numbers itself and is tested on it in each language.

This sounds like a detail and is not. A question about volume whose answer is captured at the wrong order of magnitude is worse than never asking — you would walk away believing you had a number.

Where this is honest about itself

Detecting both languages, keeping the panel and the suggestion in different languages, and converting numbers all work today.

The list of phrasings that carry evasion in Mandarin has been built — including the deferrals that sound like agreement and are not — but a bilingual business analyst still needs to check it against real client recordings before anyone should lean on it. Building the list was the long part; checking it needs a native speaker who does this work.

One thing worth knowing: the two languages deliberately do not cover the same ground. Mandarin catches a category of deferred commitment that English carries in tone rather than in words, which no word list could find. Each language is stronger than the other somewhere, and those gaps are named rather than assumed away.

When the Connection Drops

Networks fail in client offices. Elicita is built so that when the part of it that needs the internet cannot be reached, the meeting does not stop — it just gets quieter.

What keeps working

The fastest and most valuable suggestions never needed the internet at all.

Spotting an unquantified word like *fast* is a matter of matching against a list prepared in advance, so it carries on exactly as before. So does your coverage tracking, and every button on the panel.

What pauses is the slower, cleverer layer: noticing that an answer contradicts something said twenty minutes ago, or that a system just mentioned appears nowhere in the briefing pack.

What you lose, and what you keep

	Still working	Paused
Vague words like "fast" or "a few"	Yes	—
Topics you have not covered yet	Yes	—
All four one-tap responses	Yes	—
Contradicting something said earlier	—	Paused
Spotting a system nobody briefed you on	—	Paused

You are told, not left to guess

1 / 4

22:00

EN

Deterministic only

The model is unreachable. Wording and coverage triggers still fire; the slow lane is paused and will catch up.

How fast is fast?

When you say the dashboard has to be fast — what does that mean in seconds?

unquantified adjective — "fast"

Asked it

Park it

What am I missing?

Go deeper

Type a question instead...

The panel says plainly that it is running on the fast path only

The panel says which mode it is in and what that means, and it is told before it shows you a single suggestion. The alternative — going quiet and letting you assume there was nothing worth asking — is the one behaviour that would genuinely mislead you.

Nothing is invented while the connection is down. Work that needed the missing piece is recorded as not done, with the reason, and it is visible afterwards rather than silently absent.

AFTER THE MEETING


Checking the Recording

During the meeting, turning speech into text is tuned for speed — a suggestion that arrives late is worthless. Afterwards it is tuned for being right, because everything in the write-up comes from this version and not from the live one.

Two transcriptions, on purpose

The recording is transcribed twice, by two different companies.

Not for redundancy — for disagreement. Where two independent services produce the same text, it is almost certainly right and needs nobody's attention. Where they differ, something was genuinely unclear, and that is exactly what a person should look at.



Where the two transcriptions disagreed

The two readings sit side by side because you are comparing them. Stacked, it becomes a memory exercise.

Disagreements are always flagged and never quietly resolved by picking a winner. The whole value is in surfacing the doubt.

Then the audio is destroyed

The moment both transcriptions and the speaker identification have finished — whether they succeeded or failed — the recording is deleted, and the deletion is recorded.

Nothing still needs it, and keeping it only widens what could be exposed if something went wrong later.

Where this is honest about itself

Running both services, comparing them, flagging every disagreement and destroying the audio all work today.

There is now a way to measure how much this pairing actually catches: given a recording and a known-correct transcript, it reports what share of the mistakes the two services would have shown you. A sample containing no mistakes reports "cannot tell" rather than a clean pass, because it is no evidence either way.

That measurement has not yet been run against real client recordings. The answer can be produced; it has not been produced.

Getting the Write-Up

The meeting produces a transcript. What you actually need is a set of documents somebody can act on, where every claim can be traced back to the moment it came from.

What comes out

The documents

- 1 The open questions**
Ranked by how much they matter to the build, so the list is a plan rather than an inventory.
- 2 A log of decisions**
What was settled, and by whom. This is the part people forget and then argue about.
- 3 A draft project brief**
What is being built, for whom, and why — assembled from what was said rather than from a template.
- 4 A follow-up email**
Covering what you heard and what you still need, for you to read and send. It is never sent for you.


DEBRIEF

Discovery 3 — integrations and volumes

Every claim below links to the moment it came from. Nothing is written without one.

PROJECT BRIEF

Northwind needs Zephyr WMS integrated with the finance ledger and the carrier portal, with same-day visibility of consignment status for operations staff.

“Zephyr WMS pushes to the ledger nightly, and ops can never see where a consignment actually is.” — Client — IT, 21:24

inferred

OPEN QUESTIONS

What does “fast” mean for the operations dashboard, in seconds?

“The dashboard just has to be fast, that is the main thing.” — Client — Ops, 08:15

stated

Is nightly ledger sync acceptable at go-live, or is real-time required?

“Zephyr WMS pushes to the ledger nightly.” — Client — IT, 21:24

inferred

Does SLA-4 mean four hours end to end, or four hours per hop?

“The SLA is four hours end to end.” — Client — Ops, 31:42

stated

DECISIONS

Carrier portal integration is in scope for phase one.

“Carrier portal has to be in the first release, that is not negotiable.” — Client — Programme, 38:07

stated

Reporting is deferred to phase two.

“Reporting can wait — we can live with exports for a while.” — Client — Programme, 39:50

stated

The write-up, with every claim showing where it came from

Said, or worked out

This is the distinction a reviewer cares about most, so it is on the face of every claim rather than buried in it.

Two kinds of claim, always labelled


	Stated	Inferred
Where it came from	The client said it	Elicita concluded it
What is attached	The quote, with who said it and when	The quote it was concluded from
What you do with it	Check the wording	Judge whether you agree

Blurring those two would make the whole document untrustworthy, so nothing can be written without a source attached. It is not a habit or a review step: a claim without a source cannot be saved at all.

Asking about the meeting in your own words

The documents answer the questions somebody anticipated. For the ones nobody did — *what did they actually say about the March deadline, draft me the paragraph about integrations* — you can just ask.

Before you have asked anything, the screen offers a few openers and says plainly where its answers come from.



DEBRIEF

Ask about this meeting

Northwind Logistics — Discovery 2

YOUR QUESTION

Ask anything about what was said

What did they say about the March deadline?

Ask

What did they say about the March deadline?

Which requirements are still only inferred?

Draft the follow-up paragraph about integrations.

The conversation before anything has been asked


DEBRIEF

Ask about this meeting

Northwind Logistics — Discovery 2

CONVERSATION

YOU

What did they actually say about the March deadline?

ELICTA

Priya Raman (Ops, 08:15) said “the pilot has to be live before the March peak — after that nobody has time to look at it.” She did not give a date, so March 31 is inferred from “before the March peak”, not stated. It is still an open question.

YOU

Draft the follow-up paragraph about integrations.

ELICTA

Suggested: “As discussed, the pilot covers the dispatch and telematics feeds only. The warehouse management system is out of scope for this phase — we agreed to revisit it once the dispatch integration is stable.” This rests on Dan Okoro (IT, 22:40) and the scope note in your engagement brief.

YOUR QUESTION

Ask anything about what was said

What did they say about the March deadline?

Ask

Asking a question about the meeting in your own words

Answers quote and attribute rather than summarise, and say plainly when something was worked out rather than said. That is the same rule the documents follow, for the same reason: an answer you cannot check is one you have to take on trust, and this is not a product that asks for trust.

If something goes wrong partway

The write-up runs as a sequence of steps. If one of them fails, the sequence stops there and says so, rather than passing invented material to the next step.

You get a partial result and an explanation, which is recoverable. You do not get a complete-looking document with a fabricated section in it, which is not.

Where this is honest about itself

The whole sequence works, and so does asking follow-up questions in your own words. Producing a complete run needs one of the outside services with enough capacity; the connection has been tested against the live service and works, but the account used for testing was rate limited, so a full run has not yet been produced end to end.

Carrying It Into the Next Meeting

Requirements are not settled in one sitting. Everything that makes a tool which follows a whole engagement better than one that forgets after each meeting lives here.

What survives

The confirmed requirements, the decisions, and — most usefully — the questions that are still open.

 **ENGAGEMENT**

Northwind Logistics

Meetings
3

Requirements confirmed
24

Still open
3

CARRIED INTO THE NEXT MEETING

What does “fast” mean for the operations dashboard, in seconds? First raised in Discovery 1	3 meetings
Is nightly ledger sync acceptable at go-live? First raised in Discovery 3	1 meeting
Who signs off the carrier portal data-sharing agreement? First raised in Discovery 2	2 meetings

The bank for the next meeting is weighted toward these. A question that has survived two meetings is the one most worth asking in the third.

MEETINGS SO FAR

- Discovery 1 — current state**
4 Aug 2026 · 3 of 8 sections covered
- Discovery 2 — process and people**
11 Aug 2026 · 6 of 8 sections covered
- Discovery 3 — integrations and volumes**
18 Aug 2026 · 7 of 8 sections covered

What carries into the next meeting

Open questions lead the screen, ahead of completed work. That ordering is the point: what you settled is behind you, and what you did not is what you are walking back in to deal with.

The question that keeps surviving

A question that has come through three meetings unanswered is flagged.

That is worth knowing before you walk in, because there are only two explanations and both matter. Either it is genuinely hard, or it is being deflected. Flagging it puts the choice of how to handle that in front of you while you can still plan for it.

The question list for the next meeting is weighted towards these. You do not start from scratch, and you are not re-asked what was already settled.

What is deliberately not kept

Only those things above are stored. The working notes produced along the way during a write-up are rebuilt from the recording rather than kept.

That is slower. It is also correct: a stored copy could go stale against the recording it came from, and a stale copy is worse than none, because nothing about it looks wrong.

RUNNING IT WELL

Judging Whether the Suggestions Are Any Good

Whether a suggestion was worth interrupting a meeting for is a judgement, not a measurement. So it is made by a person, on real meetings — and it decides whether a new version is allowed to ship.

This is for whoever is responsible for the tool across engagements, rather than for the person running a meeting.

Replay a real meeting

A recorded meeting is run through Elicita again, exactly as it happened. Nothing is sent anywhere and nothing costs anything: it is replaying its own past behaviour, which also means the result is repeatable.

You then rate what it suggested, on two separate scales.

Two ratings, deliberately not opposites

	Useful	Embarrassing
The question it answers	Was this worth the interruption?	Would I have minded the client seeing it?
A suggestion can be	Useless without being humiliating	Humiliating regardless of how apt it was
What a bad score costs	A wasted moment	Trust you do not get back

They are kept apart on purpose. Combining them into a single score would let a high usefulness rating average an embarrassment away, and an embarrassment is not the kind of thing that should be averaged.

 **REPLAY**

Discovery 3 — replay 4

Useful when surfaced
82%
M1 needs 70%

Embarrassing
0
M2 needs zero

SUGGESTIONS

When you say fast, what does that mean in seconds? 08:19 · EN	Useful	Embarrassing
How many consignments a day at peak? 14:02 · EN	Useful	Embarrassing
Could you expand on that? 19:38 · EN	Useful	Embarrassing
Is the nightly sync a constraint or a choice? 21:31 · EN	Useful	Embarrassing

Rate embarrassing only for a suggestion you would not want a client to have seen. One is enough to hold a release.

A version that clears both bars

Two bars, and one of them is absolute

At least **70 per cent** of suggestions must be rated useful.

The number rated embarrassing must be **zero**. Not low — zero. One is enough to hold a release.


REPLAY

Discovery 3 — replay 2

Useful when surfaced
64%
M1 needs 70%

Embarrassing
1
M2 needs zero

SUGGESTIONS

When you say fast, what does that mean in seconds? 08:19 · EN	Useful	Embarrassing
How many consignments a day at peak? 14:02 · EN	Useful	Embarrassing
Could you expand on that? 19:38 · EN	Useful	Embarrassing
Has the client considered whether this project is viable at all? 26:10 · EN	Useful	Embarrassing

Rate embarrassing only for a suggestion you would not want a client to have seen. One is enough to hold a release.

A version that fails both

A version that fails either bar is blocked automatically. Nobody has to remember to check.

Where this is honest about itself

Replaying, rating, both bars and the automatic block all work today.

Both bars are judged separately for each language, so adding a language means finding a fluent rater who also does business analysis — not just translating a word list. That is worth planning for early, because it is a hiring problem rather than a software one.

What It Will Not Do

It is easier to trust a tool when you know where its edges are. This is the list of things Elicita deliberately does not do, and why each line was drawn.

It never speaks to the client

It has no voice in the meeting. It does not send messages, and it does not act on anybody's behalf. Every suggestion is shown to you, and you decide.

If you ignore it for ninety minutes, nothing happens. That is a supported way to use it, not a failure.

It is not a transcription product

It does turn speech into text, because it cannot work otherwise. But that part is bought in, and it is not what Elicita is for. If all you want is a recording and a transcript, this is a great deal of machinery for something simpler tools already do well.

It is not a note-taker

Notes come out of it, as a by-product. The purpose is the question asked in the room. A tool that captured the meeting perfectly and never improved what was asked would have missed the point entirely.

It is not something you talk to during a meeting

You cannot ask it things while the meeting is running. There is no conversation with it. It offers; you accept or decline.

That restriction is deliberate. Anything that invites you to start typing has taken your attention away from the client, which is the exact problem it was built to relieve. Asking it things in your own words is something you do afterwards, on a different screen.

It stays quiet when it is unsure

Given the choice between a suggestion that might be useful and saying nothing, it says nothing.

One suggestion that embarrasses you in front of a client does more damage than ten good ones repair. So the bar for speaking is set high, and it is checked against recordings of real meetings before every release.

The practical effect is that Elicita will sometimes miss something a sharp analyst would have caught. That trade was made knowingly, and it is the right way round.

It does not replace judgement

It does not decide what matters, what is worth pushing on, or when to let something go because the room has had enough.

It notices, it suggests, and it writes down what it heard. Everything that requires knowing this client, this project and this moment stays with you.

What Works Today

This page is assembled automatically every time this guide is made, from notes kept alongside the product itself, so it cannot fall out of date. Editing it by hand has no effect — the next build puts it back.

36 notes in total, 30 of them describing something that works today.

This page is read from the notes kept alongside each part of the product, so it says what is true now rather than what was true when somebody last remembered to update a page.

Working today

You can rely on these.

Part of the product	What the note says
Prepare for the engagement	Setting up the engagement, tagging and indexing documents, the client vocabulary list, and reviewing and pruning the question bank
Start the meeting	The consent gate, the record of who confirmed, and the whole audio-handling policy including automatic destruction
Start the meeting	The desktop app now opens a real microphone. It offers the audio interface and the silent-join capture path, warns you off the room mic, and releases the device when you stop
Start the meeting	Elicita will join meetings as a silent participant through a single meeting-bot service that covers every platform, rather than integrating with Zoom, Teams and Meet separately. That keeps one integration instead of three, and gives us who-said-what as a fact the platform reports rather than something guessed from the audio. The connection point is built and tested; the vendor contract is a commercial step, not an engineering one
Catch a vague answer while it still matters	The whole panel, all four responses, the running coverage count, and the live connection that feeds suggestions to it
Catch a vague answer while it still matters	Your taps are sent back, so the debrief knows which suggestions you used and which you set aside
Catch a vague answer while it still matters	The microphone is connected. The desktop app opens the device, and audio flows through to the suggestions on this screen
Run a meeting in two languages	Detecting and displaying both languages, keeping the suggestion and the interface in different languages, and number conversion
Run a meeting in two languages	The Mandarin list now covers the phrasings that carry evasion in Chinese business speech — including the deferrals that sound like agreement and are not, such as “we’ll look into it”, and the qualifiers that quietly withdraw most of a yes
Run a meeting in two languages	We measured it rather than leaving it open. The two languages do not cover the same ground, and that is correct: Mandarin catches a whole category of deferred commitment that English carries in tone rather than words, so no word list could find it. Each language is stronger than the other somewhere, and the gaps are now named rather than assumed away
When the connection drops	The two-speed design, the panel’s degraded state, and honest recording of what could not be completed

Part of the product	What the note says
When the connection drops	The panel is now told which mode it is in by the live connection itself, before it shows you a single suggestion — so a quiet panel is never ambiguous between “nothing to say” and “the model is unreachable”
Check the recording	Running both transcribers, comparing them, flagging every disagreement, and the automatic destruction of the audio
Check the recording	The check for that is now built: given a recording and a correct transcript, it reports what share of the mistakes the pairing would actually have shown you. A sample containing no mistakes reports “cannot tell” rather than a clean pass, because it is no evidence either way
Get the write-up	The whole sequence — cleaning up the transcript, translating while keeping the original, sorting statements into template sections, drafting the documents, and attaching every source
Get the write-up	If any step fails, the sequence stops there and says so, rather than passing invented material to the next step
Get the write-up	Asking follow-up questions in your own words now has a screen. You can query the meeting, ask for a paragraph to be drafted, and every answer names the moment it rests on
Carry what you learned into the next meeting	Carrying requirements, decisions and open questions across meetings, and weighting the next meeting's questions toward what is unresolved
Carry what you learned into the next meeting	All of it now survives restarting the service. The client, its meetings, the open questions, the standing requirements and the prepared question bank are written to a database as they change, and read back when Elicita next starts
Carry what you learned into the next meeting	Only the five things above are stored. The working notes a debrief produces along the way are rebuilt from the recording rather than kept — slower, but a stored copy could go stale against the recording it came from, and a stale one is worse than none
Judge whether the suggestions are any good	Replaying meetings, rating suggestions, both quality bars, and automatic blocking of a release that fails either
Judge whether the suggestions are any good	The starting weights are now chosen rather than left at a placeholder, and the reasoning is written down: say less, more confidently, because a suggestion that misfires in front of a client costs more than one that never fires. A question you have already asked can never be suggested again ahead of one you have not — that is guaranteed by the numbers themselves, not left to chance
Set up the services Elicita uses	Everything on this screen, including live credential testing for the AI provider and both supported transcription services
Set up the services Elicita uses	For a shared deployment the key can now come from your own secret store — 1Password, Vault, AWS, Google, anything with a command line. If the store cannot be reached, Elicita refuses to start and says why, rather than quietly making a new key that leaves the other machines unable to read anything
Control the recording while the meeting runs	The whole screen — pause and resume, the state display, the input warning, and voice enrolment
Control the recording while the meeting runs	Pause now pauses a real microphone. Audio captured while paused is discarded rather than held back, and resuming is instant because the device is never closed

Part of the product	What the note says
Get Elicita onto people's machines	Packaging for both platforms, signing and verification, the managed-deployment profile, and the automatic check that both platforms behave identically
Get Elicita onto people's machines	The About screen reports the real version and platform it is running on
Get Elicita onto people's machines	Whether the build is signed, and how it was installed, are now read from the operating system rather than taken on trust. Where the answer cannot be established the screen says so, which reads differently from "unsigned" — an IT reviewer needs to tell those apart
Get Elicita onto people's machines	Elicita checks for a new version on launch and tells you what is waiting. It never installs on its own — an update that restarted the app by itself would eventually do it during a client meeting

Needs setting up first

These work, once somebody has configured the outside service they depend on.

Part of the product	What the note says
Prepare for the engagement	Drafting the questions needs an AI provider configured in Settings — a one-off step covered in journey 10. Until then the screen works and the bank stays empty

Still to come

These are described elsewhere in this guide as things that are not finished. They are listed together here so that nobody has to hunt for the caveats.

Part of the product	What the note says
Run a meeting in two languages	A bilingual business analyst still needs to review that list against real client recordings before we would rely on it. Building it was the long part; checking it is the part we cannot do without a native speaker who does this work
Check the recording	The check still has to be run against real client recordings. We can measure the answer now; we do not have it yet
Get the write-up	Needs an AI provider with capacity. The connection has been tested against the live service and works; the account used for testing was rate limited, so a complete run has not been produced yet
Judge whether the suggestions are any good	Both bars are judged per language, so adding a language means finding a fluent rater with business-analysis experience — not just translating a word list. Worth planning for early
Get Elicita onto people's machines	The update service needs its signing key and address set before a release goes out. Until then the screen says there is no update channel configured, rather than accepting whatever it is offered